

4741 Bias and Variance
HW 4, Question 4
Aman jain (aj644@cornell.edu)

a)

```
In [1]: using Plots
        using Statistics
```

Suppose we have a sinusoid function

```
In [2]: f(x) = 10*sin.(x)
```

```
Out[2]: f (generic function with 1 method)
```

Our dataset \mathcal{D} will consist of seven datapoints drawn from the following probabilistic model.

For each datapoint we randomly draw x_i uniformly in $[0,6]$ and observe a noisy $y_i = f(x_i) + \epsilon_i$, where ϵ_i is some noise drawn from a standard normal distribution $\mathcal{N}(0, 1)$.

Generate a sample dataset from this distribution.

```
In [3]: using Random
        d = 6*rand()
        d
```

```
Out[3]: 0.6742295554456206
```

```
In [4]: h = [[1,3],[4,5]]
        h[1][1]
```

```
Out[4]: 1
```

```
In [5]: n = 7
D = zeros(n, 2)
for i = 1:n
    D[i,1] = 6*rand()
    D[i,2] = f(D[i,1])+randn() # yi = f(xi) + ei
end
D
```

```
Out[5]: 7×2 Array{Float64,2}:
 5.78846  -5.38629
 5.74309  -5.16613
 4.74145  -9.06159
 1.45705  10.3888
 4.32511  -11.0167
 3.87209  -7.30246
 2.31469   7.70446
```

Plot the dataset D and the true function $f(x)$.

```
In [6]: """plot function y=f(x)"""
function plotfunc_first(f;
    xmin=0,xmax=6,nsamples=1000)
    xsamples = range(xmin,stop=xmax,length=nsamples)
    plot(xsamples, [f(x) for x in xsamples], color="black") ## only dif
end
```

```
Out[6]: plotfunc_first
```

```
In [7]: """plot function y=f(x)"""
function plotfunc(f;
    xmin=0,xmax=6,nsamples=1000)
    xsamples = range(xmin,stop=xmax,length=nsamples)
    plot!(xsamples, [f(x) for x in xsamples], color="black")
end
```

```
Out[7]: plotfunc
```

b)

Fit a linear model to D

```
In [8]: X = [D[:,1] ones(7)]  
        y = D[:,2]  
        w = X\y
```

```
Out[8]: 2-element Array{Float64,1}:  
        -4.11675745451897  
        13.775034178463931
```

```
In [9]: function l(x)  
        y_pred = [x 1]w  
        return( y_pred[1])  
end
```

```
Out[9]: l (generic function with 1 method)
```


```
In [10]: l(0.5)
```

```
Out[10]: 11.716655451204446
```

Plot the linear model $l(x)$ together with \mathcal{D} and $f(x)$. Feel free to use our method `plotfunc(f)`.

```
In [11]: plotfunc_first(1)
          plotfunc(f)
          scatter!(D[:,1],D[:,2])
```

Out[11]:



c)

Fit a cubic model $c(x)$ to \mathcal{D}

```
In [12]: # first, construct a Vandermonde matrix
max_order = 3
x = D[:,1]
V = zeros(n, max_order+1)
for k=0:max_order
    V[:,k+1] = x.^k
end

# solve least squares problem
w_c = V\D[:,2]
w_c
```

```
Out[12]: 4-element Array{Float64,1}:
 -4.21835236453993
 22.981019642056605
 -10.378989661685075
  1.1040788162896285
```

```
In [13]: function c(x; order = max_order, w = w_c)
           y = 0
           for k=0:order
               y += w[k+1]*x^k
           end
           return y
       end
```

```
Out[13]: c (generic function with 1 method)
```


```
In [14]: c(0.5)
```

```
Out[14]: 4.815419893103307
```

Plot the cubic model with D and $f(x)$

```
In [15]: plotfunc_first(c)
          plotfunc(f)
          scatter!(D[:,1],D[:,2])
```

Out[15]:



d)

Repeat the parts (b) and (c) for 1000 different sets \mathcal{D} . Compute \bar{l} and \bar{c} , the average linear and average cubic models. (Please name \bar{l} as `l_avg(x)` and \bar{c} as `c_avg(x)` in following codes)

```

In [16]: max_iter = 1000
w_l_bar_all_iter = zeros((2, max_iter))
for j = 1:max_iter
    D_l_bar = zeros((n, 2))
    for i = 1:n
        D_l_bar[i,1] = 6*rand()
        D_l_bar[i,2] = f(D_l_bar[i,1])+randn() #  $y_i = f(x_i) + \epsilon_i$ 
    end
    X_l_bar = [D_l_bar[:,1] ones(7)]
    y_l_bar = D_l_bar[:,2]
    w_l_iter = X_l_bar \ y_l_bar

    w_l_bar_all_iter[:,j] = w_l_iter
end

w_l_bar_all_iter

```

```

Out[16]: 2×1000 Array{Float64,2}:
 -5.54828  -4.94422  -1.68622  -5.32331  ...  -4.4278  -4.01315  -4.7377
 8
 18.2672   16.1442    2.00356  19.207      13.8299  11.1007  14.052

```

```

In [17]: w_l_avg = mean(w_l_bar_all_iter, dims=2)

```

```

Out[17]: 2×1 Array{Float64,2}:
 -3.6639131625395844
 11.162068365196964

```

```

In [18]: function l_avg(x)
    y_pred = [x 1]w_l_avg
    return( y_pred[1])
end

```

```

Out[18]: l_avg (generic function with 1 method)

```

```

In [19]: l_avg(0.5)

```

```

Out[19]: 9.330111783927173

```

```

In [20]: max_iter = 1000
max_order = 3
w_c_bar_all_iter = zeros((4, max_iter))
n = 7

for j = 1:max_iter
    D_c_bar = zeros((n, 2))
    for i = 1:n
        D_c_bar[i,1] = 6*rand()
        D_c_bar[i,2] = f(D_c_bar[i,1])+randn() #  $y_i = f(x_i) + \epsilon_i$ 
    end

    x = D_c_bar[:,1]
    #print(x)
    V = zeros(n, max_order+1)
    for k=0:max_order
        V[:,k+1] = x.^k
    end

    # solve least squares problem
    w_c_iter = V\D_c_bar[:,2]
    #print(w_c_iter)
    #w_c_iter

    w_c_bar_all_iter[:,j] = w_c_iter
end

w_c_bar_all_iter

```

```

Out[20]: 4×1000 Array{Float64,2}:
-4.08815  -2.87338  -3.32336  ...  -2.56673    4.58324   -6.82033
22.5895   22.4843   20.152    20.3015   11.9309   26.6007
-9.6757  -10.2223  -9.36472  -9.237    -7.0215  -11.9851
0.993027  1.07941  1.0218    0.985538  0.813287  1.33425

```

```

In [21]: w_c_avg = mean(w_c_bar_all_iter, dims=2)

```

```

Out[21]: 4×1 Array{Float64,2}:
-2.9469356018885002
20.817732508578587
-9.434365022562043
0.9902613312517625

```

```

In [22]: c_avg(x) = c(x; order = max_order, w = w_c_avg)

```

```

Out[22]: c_avg (generic function with 1 method)

```




```
In [23]: c_avg(0.5)
```

```
Out[23]: 5.227122063166754
```

Plot \bar{l} together with $f(x)$.

```
In [24]: plotfunc_first(l_avg)
plotfunc(f)
```


```
Out[24]:
```



Plot $\bar{c}(x)$ together with $f(x)$.

```
In [25]: plotfunc_first(c_avg)
plotfunc(f)
```

Out[25]:



e)

Compute the bias of \bar{l} . You can use our integrate function.

```
In [26]: function integrate(f, a, b)
    n = 1000
    delta = (b - a)/n;                ## nothing to change below here
    xs = a*ones(n) + [0:1:n-1;] * delta;    ## n, right is 1:n * d
    fx = map(f, xs);
    return sum(fx) * delta
end
```

Out[26]: integrate (generic function with 1 method)

```
In [27]: g_l(x) = (f(x) - l_avg(x))^2
```


Out[27]: g_l (generic function with 1 method)

```
In [28]: bias_l_bar = integrate(g_l,0,6)
```

```
Out[28]: 103.88269898743333
```

```
In [29]: x_range = range(0,stop=6,length=1000)
x = [x for x in x_range]
y = f(x)
y_2 = [l_avg(x) for x in x_range]
plot(x, y, color="red", linewidth=2.0)
plot!(x, y_2, fillrange=[y y_2], color="black", linewidth=2.0)
```

```
Out[29]:
```



Compute the bias of the cubic model $\bar{c}(x)$.

```
In [30]: g_c(x) = (f(x) - c_avg(x))^2
```

```
Out[30]: g_c (generic function with 1 method)
```


```
In [31]: bias_c_bar = integrate(g_c,0,6)
```

```
Out[31]: 5.355579462582125
```

We can interpret the bias as how far off our averaged model is from the true function. One way to visually see this is by plotting $\bar{l}(x)$ with $f(x)$ and color in their difference. Try out the plotting function below.

```
In [32]: x_range = range(0,stop=6,length=1000)
x = [x for x in x_range]
y = f(x)
y_2 = [c_avg(x) for x in x_range]
plot(x, y, color="red", linewidth=2.0)
plot!(x, y_2, fillrange=[y y_2], color="black", linewidth=2.0)
#plot!([x y_2], fillrange=[y y_2], fillalpha=0.3, c=:orange)
```

Out[32]:



Which model has smaller bias?

Polynomial model has a smaller bias (of 4.6) as compared to Linear model (which has a bias of 104).

f)

Next compute the variance of the linear model.

```
In [33]: x_range = range(0,stop=6,length=1000)
x = [x for x in x_range]
Var_l = 0

for i = 1:1000
    #print(x[i])
    X_l = [x[i] 1]
    #print(i, X_l)
    #print(i,w_l_bar_all_iter[:,i])
    y_jth_model_pred = X_l*w_l_bar_all_iter[:,i]
    #print(y_jth_model_pred)
    #print(w_l_avg)
    y_l_avg_pred = X_l*w_l_avg
    Var_l = Var_l + (y_jth_model_pred[1] - y_l_avg_pred[1])^2
end

variance_linear = Var_l*6/1000 #integration factor (b-a/n)
```

Out[33]: 50.37938888279812

Compute the variance of the cubic model.

```
In [34]: x_range = range(0,stop=6,length=1000)
x = [x for x in x_range]
Var_c = 0

for i = 1:1000
    #print(x[i])
    X_l = x[i]
    #print(i, X_l)
    #print(i,w_l_bar_all_iter[:,i])
    y_jth_model_pred = c(X_l, w=w_c_bar_all_iter[:,i])
    #print(y_jth_model_pred)
    #print(w_l_avg)
    y_c_avg_pred = c(X_l, w=w_c_avg)
    Var_c = Var_c + (y_jth_model_pred[1] - y_c_avg_pred[1])^2
end

variance_cubic = Var_c*6/1000
```

Out[34]: 426.55981691979196

```
In [35]: total_error_linear = variance_linear + bias_l_bar  ## I am doing square wh
total_error_cubic = variance_cubic + bias_c_bar
print("Total linear error: ",total_error_linear,"Total cubic error: ",total

Total linear error: 154.26208787023145Total cubic error: 431.915396382
37407
```

Which model had higher variance? How do you interpret this? Which model has smaller overall error?

**Cubic model has higher variance (426) as compared to linear model (50).
Total error is high for cubic model (431) as compared to linear models (154).**

This indicates typical bias-variance tradeoff. Models with high values of bias (linear model in this case), undermine their training dataset (underfit) leading to not learn the trends in training dataset. consequently these models are simpler also. These models do not vary a great deal with change in training dataset.

High-variance are more complex and represent their training set well (at times overfit). However, these models are not able to handle random variation in test dataset and hence do not generalize well.

Ideally, one wants to choose a model with training error being low and generalize on out of set data as well - Alas! not easy to get there.

g)

How do you think your results would depend on the number of points in the data set D ? Feel free to perform an experiment to check. How many points would you need before the opposite model has smaller overall error?

```
In [36]: n_min = 7 # min data points
n_max = 50 #max data points
max_iter = 1000
x_range = range(0,stop=6,length=max_iter)
x = [x for x in x_range]
n_iter = n_max-n_min+1
max_order = 3
```

```

Variance_matrix = zeros((2,n_iter))
Bias_matrix = zeros((2,n_iter))
total_error_matrix = zeros((2,n_iter))

for n = n_min:n_max

    ## Linear model calculations start here
    w_l_bar_all_iter = zeros((2, max_iter))
    w_c_bar_all_iter = zeros((4, max_iter))

    for j = 1:max_iter
        D_bar = zeros((n, 2))
        for i = 1:n
            D_bar[i,1] = 6*rand()
            D_bar[i,2] = f(D_bar[i,1])+randn() #  $y_i = f(x_i) + \epsilon_i$ 
        end
        X_l_bar = [D_bar(:,1) ones(n)]
        y_l_bar = D_bar(:,2)
        w_l_iter = X_l_bar \ y_l_bar
        w_l_bar_all_iter(:,j) = w_l_iter

        b = D_bar(:,1)
        V = zeros(n, max_order+1)
        for k=0:max_order
            V(:,k+1] = b.^k
        end
        w_c_iter = V \ D_bar(:,2]
        w_c_bar_all_iter(:,j) = w_c_iter

    end

    w_l_avg = mean(w_l_bar_all_iter, dims=2)
    Var_l = 0
    Bias_l = 0

    w_c_avg = mean(w_c_bar_all_iter, dims=2)
    Var_c = 0
    Bias_c = 0

    for k = 1:max_iter
        X_l = [x[k] 1]
        y_jth_model_pred = X_l*w_l_bar_all_iter(:,k]
        y_l_avg_pred = X_l*w_l_avg
        Var_l = Var_l + (y_jth_model_pred[1] - y_l_avg_pred[1])^2
        Bias_l = Bias_l + (f(x[k][1])-y_l_avg_pred[1])^2

        X_c = x[k]
        y_jth_model_pred = c(X_c, w=w_c_bar_all_iter(:,k])
        y_c_avg_pred = c(X_c, w=w_c_avg)
    end
end

```

```

Var_c = Var_c + (y_jth_model_pred[1] - y_c_avg_pred[1])**2
Bias_c = Bias_c + (f(x[k][1]) - y_c_avg_pred[1])**2

end
Variance_matrix[1,n-n_min+1] = Var_l*(6/max_iter)    ##integration fa
Bias_matrix[1,n-n_min+1] = Bias_l*(6/max_iter)        ##integration fa

Variance_matrix[2,n-n_min+1] = Var_c*(6/max_iter)
Bias_matrix[2,n-n_min+1] = Bias_c*(6/max_iter)

end

total_error_matrix = Variance_matrix + Bias_matrix


#print(Variance_matrix)
#print(Bias_matrix)

print(total_error_matrix[:,1:5])
plot(total_error_matrix[1,:])
plot!(total_error_matrix[2,:])

[159.11809730300524 152.6357284653181 145.24229994630056 136.516953270
8011 131.68946169239317; 442.33623019159825 116.32572754368468 141.147
0892034337 54.70764400947264 23.785039529958958]

```

Out[36]:



As number of datapoints in D increases, bias comes down significantly. As $n > 9$, opposite model (in this case Cubic model) has lower total error than part (f). Errors for both models flattens out after 30 odd data points.

h)

Instead of sampling new data to compute the bias and variance of our model, we could use a bootstrap estimator to get more use out of the few data points we have. Try this for a few different data set sizes and report on your results. How big a data set is needed for the bootstrap to give a reliable estimate of the bias and variance?

Plan:

1. In order to understand how many data points are needed for a reliable measure of true error (which in turn will give reliable measures for bias and variance), I am going to take $n = 7$ and plot mean error square $((y_{\text{pred}} - y_{\text{true}})^2)$ for different values of K (bootstrap sample size). At some point of K , this error should flatten out - which will give us a threshold for K .
2. Try exercise 1 for couple of different values of n .
3. Understanding the affect of n and K on Confidence intervals for w .

Note: I will be restricting this exercise of bootstrapping to linear models and look at variation of first coefficient of w in those models. Conclusions from this should hold for other models and other values of w as well.

Part 1

Mean error vs K (for $n = 7$)

```
In [107]: n_h_1 = 7

## Generate a sample size of 500 points from which we will be drawing samples
Sample_size = 500
Samples = zeros((Sample_size, 2))
for i = 1:Sample_size
    Samples[i,1] = 6*rand()
    Samples[i,2] = f(Samples[i,1])+randn() # yi = f(xi) + ei
end

Boot_Strap_Sample_Size_max = 1000
Avg total Bootstrap error = zeros(Boot Strap Sample Size max)
```


```


for m = 1:Boot_Strap_Sample_Size_max
    loop_error = 0
    for k = 1:m
        error = 0
        D = zeros((n_h_1,2))
        for j = 1:n_h_1
            u = rand(1:Sample_size)
            D_random_point = Samples[u,:]
            #print(D_random_point)
            D[j,:] = D_random_point
        end
        #print(D)
        X_l_bar = [D[:,1] ones(n_h_1)]
        y_l_bar = D[:,2]
        w_l_iter = X_l_bar \ y_l_bar
        y_pred = X_l_bar * w_l_avg
        #print(y_pred)
        error = sum((y_pred - y_l_bar).^2)
        loop_error = loop_error + error
    end
    Avg_total_Bootstrap_error[m] = loop_error/m
end

plot(Avg_total_Bootstrap_error)
#true_error_plot = true_error*ones(size(boot_error))
#plot!(true_error_plot)
#size(true_error)
#print(true_error)

```

Out[107]:





As we see above for about $K = 250$, errors reduce in a constant range.

Part 2

Part 1 for $n = 25, 45, 60$.

```
In [109]: n_h_1 = 25


## Generate a sample size of 500 points from which we will be drawing sa
Sample_size = 500
Samples = zeros((Sample_size, 2))
for i = 1:Sample_size
    Samples[i,1] = 6*rand()
    Samples[i,2] = f(Samples[i,1])+randn() #  $y_i = f(x_i) + \epsilon_i$ 
end

Boot_Strap_Sample_Size_max = 1000
Avg_total_Bootstrap_error = zeros(Boot_Strap_Sample_Size_max)
for m = 1:Boot_Strap_Sample_Size_max
    loop_error = 0
    for k = 1:m
        error = 0
        D = zeros((n_h_1,2))
        for j = 1:n_h_1
            u = rand(1:Sample_size)
            D_random_point = Samples[u,:]
            #print(D_random_point)
            D[j,:] = D_random_point
        end
        #print(D)
        X_l_bar = [D[:,1] ones(n_h_1)]
        y_l_bar = D[:,2]
        w_l_iter = X_l_bar \ y_l_bar
        y_pred = X_l_bar * w_l_avg
        #print(y_pred)
        error = sum((y_pred - y_l_bar).^2)
        loop_error = loop_error + error
    end
    Avg_total_Bootstrap_error[m] = loop_error/m
end

plot(Avg_total_Bootstrap_error)
#true_error_plot = true_error*ones(size(boot_error))
#plot!(true_error_plot)
#size(true_error)
```

```
#print(true_error)
```

Out[109]:



```
In [111]: n_h_1 = 45

## Generate a sample size of 500 points from which we will be drawing sa
Sample_size = 500
Samples = zeros((Sample_size, 2))
for i = 1:Sample_size
    Samples[i,1] = 6*rand()
    Samples[i,2] = f(Samples[i,1])+randn() # yi = f(xi) + ei
end

Boot_Strap_Sample_Size_max = 1000
Avg_total_Bootstrap_error = zeros(Boot_Strap_Sample_Size_max)
for m = 1:Boot_Strap_Sample_Size_max
    loop_error = 0
    for k = 1:m
        error = 0
        D = zeros((n_h_1,2))
        for j = 1:n_h_1
            u = rand(1:Sample_size)
            D_random_point = Samples[u,:]
            #print(D_random_point)
            D[j,:] = D_random_point
        end
        #print(D)
    end
end
```


```

X_l_bar = [D[:,1] ones(n_h_1)]
y_l_bar = D[:,2]
w_l_iter = X_l_bar \ y_l_bar
y_pred = X_l_bar * w_l_avg
#print(y_pred)
error = sum((y_pred - y_l_bar).^2)
loop_error = loop_error + error
end
Avg_total_Bootstrap_error[m] = loop_error/m
end

plot(Avg_total_Bootstrap_error)
#true_error_plot = true_error*ones(size(boot_error))
#plot!(true_error_plot)
#size(true_error)
#print(true_error)

```

Out[111]:



In [112]: n_h_1 = 70

```

## Generate a sample size of 500 points from which we will be drawing sa
Sample_size = 500
Samples = zeros((Sample_size, 2))
for i = 1:Sample_size
    Samples[i,1] = 6*rand()
    Samples[i,2] = f(Samples[i,1]) + randn() #  $y_i = f(x_i) + \epsilon_i$ 
end

```


```


Boot_Strap_Sample_Size_max = 1000
Avg_total_Bootstrap_error = zeros(Boot_Strap_Sample_Size_max)
for m = 1:Boot_Strap_Sample_Size_max
    loop_error = 0
    for k = 1:m
        error = 0
        D = zeros((n_h_1,2))
        for j = 1:n_h_1
            u = rand(1:Sample_size)
            D_random_point = Samples[u,:]
            #print(D_random_point)
            D[j,:] = D_random_point
        end
        #print(D)
        X_l_bar = [D(:,1) ones(n_h_1)]
        y_l_bar = D(:,2)
        w_l_iter = X_l_bar\y_l_bar
        y_pred = X_l_bar*w_l_avg
        #print(y_pred)
        error = sum((y_pred-y_l_bar).^2)
        loop_error = loop_error + error
    end
    Avg_total_Bootstrap_error[m] = loop_error/m
end

plot(Avg_total_Bootstrap_error)
#true_error_plot = true_error*ones(size(boot_error))
#plot!(true_error_plot)
#size(true_error)
#print(true_error)

```

Out[112]:





In []:

Part 3

We can borrow some learnings from part g, where we saw affect of size of dataset (n) on total error. Total errors flattens after a certain point (somewhere around $n = 27$).

We can further study impact of n on bootstrapping by taking three values of n . $n = 7, 15, 25$. Also I will be considering number of datapoints taken, $k = 100$. Below we will be looking at the Confidence interval for each value of n . Half width for 95% Confidence interval should keep reducing as n increases.

```
In [37]: n_h = 35
K = 1000
max_iter = K
w_l_bar_all_iter = zeros((2, max_iter))
for j = 1:max_iter
    D_l_bar = zeros((n_h, 2))
    for i = 1:n_h
        D_l_bar[i,1] = 6*rand()
        D_l_bar[i,2] = f(D_l_bar[i,1])+randn() # yi = f(xi) + ei
    end
    X_l_bar = [D_l_bar[:,1] ones(n_h)]
    y_l_bar = D_l_bar[:,2]
    w_l_iter = X_l_bar\y_l_bar

    w_l_bar_all_iter[:,j] = w_l_iter
end

z = w_l_bar_all_iter[1,:] #first coefficent of w
UB = mean(z) + 1.96*sqrt(var(z)/K) # 95% confidence interval
LB = mean(z) - 1.96*sqrt(var(z)/K)
print("n = ", n_h, "; k = ", K, "; Confidence interval = ", LB, ", ", UB,
n = 35; k = 1000; Confidence interval = -3.518638495232864,-3.45718432
75311056; Half width = 0.0307270838508793
```

n = 7; k = 100; Confidence interval = -3.8342240783025403,-3.301831389430105; Half width = 0.26619634443621765

n = 15; k = 100; Confidence interval = -3.807226401341788,-3.4520754000707012; Half width = 0.17757550063554328

n = 25; k = 100; Confidence interval = -3.6357753693738997,-3.4047754358186473; Half width = 0.1154999667776262

n = 35; k = 100; Confidence interval = -3.5271356160037763,-3.3509984924327982; Half width = 0.08806856178548905

n = 100; k = 100; Confidence interval = -3.524302390585087,-3.4115248830183442; Half width = 0.05638875378337138

n = 35; k = 10; Confidence interval = -3.7530071105002247,-3.0721762268779265; Half width = 0.34041544181114913

n = 35; k = 50; Confidence interval = -3.588653833079747,-3.2770852642801134; Half width = 0.15578428439981673

n = 35; k = 100; Confidence interval = -3.6330291669220243,-3.4462648808811265; Half width = 0.0933821430204489

n = 35; k = 1000; Confidence interval = -3.5192832189851346,-3.457504226763966; Half width = 0.030889496110584336

As we can see from above data points:

- 1. As n and k increases half width reduces, which means with more confidence we can locate true value of $w[1]$.**
- 2. this makes intuitive sense, as higher the K (number of samples withdrawn), more close the estimates will be to actual.**
- 3. For our use case n = 25-35 and K~100 should give 95% confidence interval of half width within 1.5% of the mean can be achieved.**

Rough Work

```
In [76]: n_h_1 = 100
true_error = zeros(n_h_1)

D_l_bar = zeros(n_h_1, 2)

for i = 1:n_h_1
    D_l_bar[i,1] = 6*rand()
    D_l_bar[i,2] = f(D_l_bar[i,1])+randn() #  $y_i = f(x_i) + \epsilon_i$ 
```



```

end
X_l_bar = [D_l_bar[:,1] ones(n_h_1)]
y_l_bar = D_l_bar[:,2]
w_l_iter = X_l_bar \ y_l_bar
y_pred = X_l_bar * w_l_iter
error = (y_pred - y_l_bar).^2
true_error = sum((y_pred - y_l_bar).^2) / n_h_1

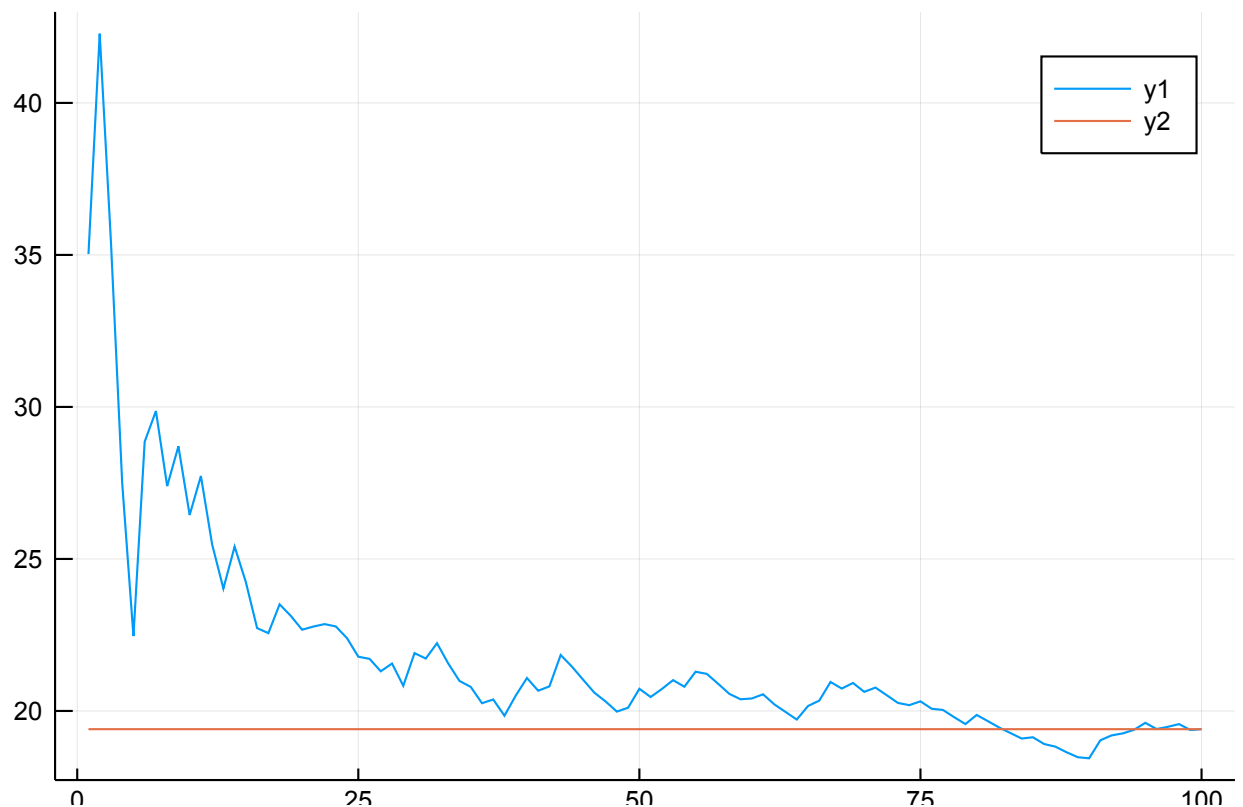
##Boot strapping
K_min = 1
K_max = n_h_1
boot_error = zeros(n_h_1)

for i = K_min:K_max
    error_i = 0
    for j = 1:i
        u = rand(1:n_h_1)
        error_i = error_i + error[j]
    end
    boot_error[i] = error_i / i
end

plot(boot_error)
true_error_plot = true_error * ones(size(boot_error))
plot!(true_error_plot)
#size(true_error)
#print(true_error)

```

Out[76]:



```

In [71]: n_h_1 = 25
         K_min = 10
         K_max = 1000
         true_error = zeros(0)

         for j = K_min:K_max
             max_iter = j
             w_l_bar_all_iter = zeros((2, max_iter))
             error = 0
             for j = 1:max_iter
                 D_l_bar = zeros((n_h_1, 2))
                 for i = 1:n_h_1
                     D_l_bar[i,1] = 6*rand()
                     D_l_bar[i,2] = f(D_l_bar[i,1])+randn() #  $y_i = f(x_i) + \epsilon_i$ 
                 end
                 X_l_bar = [D_l_bar(:,1) ones(n_h_1)]
                 y_l_bar = D_l_bar(:,2)
                 w_l_iter = X_l_bar \ y_l_bar
                 y_pred = X_l_bar*w_l_iter
                 error = error + sum((y_pred-y_l_bar).^2)
                 #print(error)
             end
             append!(true_error, error/j)
         end

         plot(true_error)
         UB = mean(true_error) + 1.67*sqrt(var(true_error)/(K_max-K_min)) # 95% c
         LB = mean(true_error) - 1.67*sqrt(var(true_error)/(K_max-K_min))
         UB_plot = UB*ones(size(true_error))
         LB_plot = LB*ones(size(true_error))
         plot!(UB_plot)
         plot!(LB_plot)
         #print("n = ", n_h, "; k = ", K, "; Confidence interval = ", LB, ", ", UB,
         #size(true_error)
         #print(true_error)

```

Out[71]:

